

Conditional diffusion emulator of Veros ACC equilibrium states: first end-to-end results

23 September 2026

Summary

- Goal: generate the post-spin-up (T, S) state of the Veros ACC channel from the EKE-closure coefficients (c_k, c_ϵ) , skipping the spin-up.
- Pipeline: DINO-Fusion diffusion code adapted to a parameter ensemble; one split runs end to end in under one A100 hour.
- Hold-out RMSE 0.061 / 0.076 / 0.141 K (scattered / band of three rows / top row), against 0.037 / 0.116 / 0.170 K for the nearest training run and 0.155 / 0.197 / 0.470 K for the mean state: the model wins in the two large-gap splits and loses only to one-step neighbours. A Wasserstein metric on the stratification profiles gives the same ranking.
- Main lesson: DINO’s constraint that re-imposes exact zeros on land and padding after every sampling step cost 0.03–0.09 K (0.152 / 0.135 / 0.173 K before). Re-imposing them at the step’s noise level, as in training, fixes it without retraining.
- Physical miss: temperature inversions are generated at the rate of the weakly eddying runs in every regime.

1 Data

- 100 Veros ACC runs on a 10×10 log-spaced grid, $c_k \in [0.0125, 0.8]$, $c_\epsilon \in [0.0875, 5.6]$; 100 years, 30-day output; interior grid $30 \times 42 \times 15$ (x, y, z).
- Kept: T and S , last 20 years (241 snapshots per run). $S = 35$ psu everywhere in every run.
- Nearly deterministic end states: within-window std 0.01 K above 1400 m vs 0.05 K between grid neighbours; the bottom level still drifts by 0.45 K over the window.
- One effective parameter: $\log c_k - 1.05 \log c_\epsilon$ explains 99 % of the variance of the domain-mean T ; the low- c_k/c_ϵ part of the grid saturates.
- Per-level distribution of T (Figure 1): above 400 m a mixture of restoring values and interior, with σ_z set by the meridional gradient, so $\mu_z \pm 3\sigma_z$ is wider than the level and the normalised values stay within $[-0.85, 0.6]$; below 1400 m a warm tail to $x' = +2.7$. Run-to-run signal σ_z^{runs} : 0.7 % of $3\sigma_z$ at the surface, 26 % at the bottom.

2 Method

- Fields $\rightarrow$ 30 channels $(T_{1..15}, S_{1..15})$, normalised per level, land $\rightarrow 0$, zero-padded $42 \times 30 \rightarrow 48 \times 32$:

$$x'_z = \frac{x_z - \mu_z}{3 \max(\sigma_z, 0.05)}. \quad (1)$$

μ_z, σ_z are computed once on all 100 runs and held fixed for every split (a deliberate, mild leakage of 30 scaling constants; no held-out state enters training).

- Denoiser: UNet2DModel (64, 64, 128, 128), 6.9 M parameters. The condition c , the standardised pair $(\log c_k, \log c_\epsilon)$, enters through an MLP added to the timestep embedding.

Distribution of T per level, all 100 runs $\times$ 241 snapshots, water cells (surface at the top). Black: μ_z ; red dashed: $\mu_z \pm 3\sigma_z$; orange ticks: the 100 per-run level means.




Figure 1: Distribution of T per level over all 100 runs and 241 snapshots (water cells), before (left) and after (right) the per-level normalisation of Eq. (1). Black: μ_z ; red dashed: $\mu_z \pm 3\sigma_z$; orange ticks: the 100 per-run level means.

- DDPM, ϵ -prediction, 1000 steps, squared-cosine schedule, EMA:

$$\mathcal{L}(\theta) = \mathbb{E}_{t, x_0, \epsilon} \left\| \epsilon - \epsilon_\theta(\sqrt{\alpha_t} x_0 + \sqrt{1 - \alpha_t} \epsilon, t, c) \right\|^2. \quad (2)$$

20 000 steps, batch 32, AdamW 10^{-4} with cosine decay; 22 min on one A100.

- Sampling: 1000-step DDPM loop with EMA weights; the predicted clean state is clipped at $|x'| \leq 1$; after every step the land and padding cells are reset to zero at the step's noise level, $\sqrt{1 - \alpha_t} z$, as in training. 32 samples per condition.

Differences from the DINO-Fusion code.

- Data: a parameter ensemble of 100 runs held in memory and split by run, instead of one DINO trajectory streamed from webdataset tars; no `tensorDict`.
- Conditioning: the UNet receives $(\log c_k, \log c_e)$ through the timestep embedding, with optional classifier-free dropout; DINO's UNet is unconditional and is steered only by sampling-time constraints.
- Transforms: the same concatenate-normalise-mask-pad chain and the same 3-std normalisation, but no vertical striding (all 15 levels kept), padding (1, 1, 3, 3) for the 42×30 grid, and statistics fixed on all runs.
- Training: the same DDPM settings, run step-based on one GPU with a cosine schedule and checkpoint resume, instead of DINO's epoch loop under `accelerate`.
- Sampling: the same DDPM loop and constraint hook; DINO's BorderZeroConstraint (exact zeros) is replaced by zeros at the step's noise level; DINO's stratification constraints are not ported yet.
- Added: hold-out splits, baselines, and the W_1 metric on stratification profiles.

3 Evaluation

- Truth: 20-year time-mean state of each held-out run; metrics on water cells.
- Baselines: mean training state; nearest training run in standardised log-parameter space; average of the axis-neighbour training runs on the grid.

- Three hold-out splits, same data file and statistics, different training sets: scattered (10 interior points), band of three rows ($c_k = 0.126, 0.2, 0.3175$; 30 runs; training rows a factor 6.3 apart), top row ($c_k = 0.8$; 10 runs; extrapolation).
- Distributional metric: each state is reduced to its profile of horizontal means over water cells, $\bar{T}_z$; per level, the Wasserstein-1 distance between the 32 generated values and the 241 true snapshots of the window, through the quantile functions,

$$W_1^{(z)} = \int_0^1 |F_{g,z}^{-1}(q) - F_{t,z}^{-1}(q)| dq, \quad W_1 = \sum_z \frac{\Delta z_z}{H} W_1^{(z)}, \quad (3)$$

with Δz_z the level thicknesses ($H = 2080$ m). All 241 snapshots are used: the quantile form accepts unequal sizes and a subsample of 32 would only add noise. A point prediction x (the baselines) gives $W_1^{(z)} = \frac{1}{241} \sum_i |x_z - \bar{T}_{z,i}|$. W_1 rewards a correct spread and does not reward collapse to the mean; the RMSE keeps the horizontal structure.

	scattered 10 runs	band of three rows 30 runs	top row 10 runs
diffusion, ensemble mean of 32: RMSE (K)	0.061	0.076	0.141
diffusion, single sample: RMSE (K)	0.147	0.160	0.208
nearest training run: RMSE (K)	0.037	0.116	0.170
neighbour average: RMSE (K)	0.020	0.116	0.170
mean training state: RMSE (K)	0.155	0.197	0.470
diffusion: domain-mean bias (K)	+0.042	+0.020	−0.016
ensemble spread / true spread (K)	0.128 / 0.031	0.132 / 0.031	0.132 / 0.032
T inversions (interfaces), generated / truth	10.2 % / 7.7 %	10.9 % / 1.3 %	12.3 % / 0.1 %
max salinity error (psu)	0.013	0.014	0.017
W_1 profile, diffusion (K)	0.066	0.075	0.132
W_1 profile, neighbour average (K)	0.036	0.095	0.142
W_1 profile, nearest training run (K)	0.044	0.095	0.142
W_1 profile, mean training state (K)	0.136	0.160	0.420
same models, exact zeros after every step: RMSE (K)	0.152	0.135	0.173
same models, exact zeros after every step: W_1 (K)	0.127	0.111	0.143

Table 1: Hold-out metrics per split, mean over the held-out runs, water cells, 32 samples per run. Last two rows: the same models sampled with DINO’s constraint (exact zeros on land and padding after every step).

4 Findings

- **The land/padding constraint was the main error source.** Same models, sampling only: exact zeros after every step give 0.152 / 0.135 / 0.173 K; zeros at the step’s noise level give 0.061 / 0.076 / 0.141 K. In training those cells carry noise like every other, so a solid block of zeros at high noise levels is out of distribution. The gain is largest at the bottom ($0.34 \rightarrow 0.07$ K, scattered); the mean-state collapse of the first version was this artefact.
- **Model vs baselines.** Band: beats the nearest run at every level below 26 m and in all three rows (0.063 / 0.074 / 0.091 K vs 0.063 / 0.132 / 0.152 K for $c_k = 0.126 / 0.2 / 0.3175$). Top row: beats the nearest run at every level but the surface and the bottom, and beats copying the row below for c_e 0.14–1.4 (0.05–0.26 vs 0.08–0.31 K); loses at the corner (0.41 vs 0.35 K) and in the flat regime ($c_e \geq 3.5$). Scattered: loses to the one-step neighbours (0.061 vs 0.037 / 0.020 K) except between 490 and 650 m (Figures 2–4).
- **Depth structure.** 0.04–0.08 K at every level in the scattered split; the bottom levels are no longer special (0.07 K) except in the top row (0.29 K, the corner runs). Worst band: the top 100 m, where $3\sigma_z \approx 12$ K amplifies sampler noise.
- **The clip is a regulariser, not a range limit.** Same models, exact zeros, scattered / top: clip at $|x'| \leq 1$ gives 0.152 / 0.173 K, at 1.5 0.194 / 0.183 K, at 3 0.302 / 0.277 K, although the final samples barely exceed $|x'| = 1$ (0.1 % of values): clipping the early clean-state estimates keeps the


chain on track. Its price is a range limit: 12% of the held-out top row's true bottom values lie above $\mu_z + 3\sigma_z$ and cannot be generated; the corner run is the largest top-row error.

- **Inversions.** 10–12 % of interfaces with $\partial_z T < 0$ in every regime; truth 7.7 % (scattered) $\rightarrow$ 1.3 % (band) $\rightarrow$ 0.1 % (top). Not learned.
- **W_1 agrees with the RMSE ranking.** Training grid points score like held-out ones (0.083 vs 0.071 K, scattered); from 490 m down the samples beat the point baselines in the band and top splits (0.058 vs 0.115 K at 650 m, band), which pay for the spread they cannot represent.
- **Plumbing.** S within 0.02 psu of 35 with no constraint; land exact; ensemble spread $4\times$ the true within-window spread (Figures 5–7).


5 Next steps

- Stratification constraint at sampling time (DINO's isotonic projection on T): targets the inversions.
- Top-row corner: the clip's range limit is now the dominant extrapolation error; a clip that follows the conditioning (bounds from the nearest training rows) is the next sampling-only test.
- Stronger conditioning: classifier-free guidance; the year as a third condition (bottom drift of 0.45 K inside the window). A second training seed to size the seed noise.
- Evaluate against snapshots as well as the time mean (nearest-snapshot RMSE, spread-skill).
- Then the same pipeline on a global Veros configuration (active S , seasonal cycle, multi-century spin-up).

Reproducibility. Branch `miniveros` of the DINO-Fusion repository, package `miniveros/`. Training at `2b542aa`; sampling and evaluation at `bfee32f` (`fill_mode` noised; the exact-zero rows: clean). Extraction 2.5 min (8 cores), training 17–22 min (one A100), generation of the hold-out and grid sets plus evaluation 14–18 min.




(a) Over the parameter grid: domain-mean T of the truth and of the generated states, their difference, and W_1 of the profile distribution. White dots: training runs; red crosses: held-out runs.




(b) RMSE (left) and W_1 (right) by depth, mean over the held-out runs.

Figure 2: Scattered split (10 interior held-out runs), 32 samples per run.



(a) Over the parameter grid: domain-mean T of the truth and of the generated states, their difference, and W_1 of the profile distribution. White dots: training runs; red crosses: held-out runs.



(b) RMSE (left) and W_1 (right) by depth, mean over the held-out runs.

Figure 4: Top row ($c_k = 0.8$ held out, extrapolation), 32 samples per run.




Figure 5: Scattered split, $c_k = 0.2$, $c_{\epsilon} = 2.222$: true 20-year mean (left) and three random samples at the surface, 650 m and the bottom; T (top three rows), S (bottom three rows). Grey: land.




Figure 6: Band-of-three split, same condition ($c_k = 0.2$, $c_{\epsilon} = 2.222$, now in the middle held-out row).




Figure 7: Top-row split, $c_k = 0.8$, $c_\epsilon = 0.5556$ (extrapolation one row above the training grid).